

LOOKS THE SAME, ANSWERS DIFFERENTLY: FLIP-DIRECTION STEERING FOR ROBUST VISION-LANGUAGE REASONING

Yeonsung Jung¹, Joonhyun Jeong^{1,2}, Hoang Pham^{3,4},
Joowon Kim¹, Yoonsik Park¹, Viet Lai⁵, Eunho Yang^{1,6}

¹Graduate School of AI, KAIST

²NAVER Cloud

³Department of Computer Science and Engineering, The Ohio State University

⁴Department of Biomedical Informatics, The Ohio State University

⁵Adobe Research ⁶AITRICS

Correspondence: ys.jung@kaist.ac.kr

ABSTRACT

Vision-language models (VLMs) are increasingly used for visual reasoning, yet subtle changes from routine image capture and processing can alter their predictions even when the images appear nearly identical. Existing inference-time defenses leverage identifiable contrasts, such as truthful versus hallucinated representations, whereas mild visual perturbations can change or preserve an answer depending on the image and prompt. We investigate these perturbation-induced prediction changes, which we call *answer flips*, through their activation shifts and decoding behavior. In the settings analyzed, a small number of directions account for a substantial fraction of clean-to-flip activation changes, and divergence tokens tend to have lower next-token margins. Based on these findings, we propose FLIPDIR (Flip-Direction Steering), a training-free inference-time method that estimates a flip-associated subspace from offline calibration pairs and selectively attenuates the hidden-state component within this subspace at uncertain decoding steps. This intervention aims to recover original predictions while preserving unaffected answers, without weight updates or a corresponding clean image at inference. We further introduce ANSWERFLIPBENCH, a reproducible benchmark spanning scientific reasoning, robot-scene understanding, and medical VQA. Across 18 settings covering two VLMs and nine combinations of datasets and perturbations, FLIPDIR improves the harmonic mean of clean preservation, non-flip preservation, and flip recovery by 2.8 points on average over the strongest baseline in each setting, with a maximum gain of 6.3 points. We will publicly release the complete benchmark, its construction code, and all experimental code.

1 INTRODUCTION

Recent vision-language models (VLMs) achieve strong performance across scientific reasoning, robot-scene understanding, and medical visual question answering (Zhang et al., 2024; Chen et al., 2025; Li et al., 2023a). These advances build on multimodal training that extends the language understanding and reasoning capabilities of large language models (LLMs) to visual inputs. However, applying these capabilities reliably requires models to extract relevant visual information and use it consistently during generation, both of which remain challenging (Covert et al., 2025; Liu et al., 2026). These limitations make robustness to visual input an important concern alongside task performance (Rosenfeld et al., 2025).

These reliability concerns have motivated several lines of work, including adversarial robustness, multimodal safety, and visual hallucination. Adversarial attacks deliberately modify images to induce incorrect or attacker-chosen outputs (Xie et al., 2025; Zhang et al., 2025). Multimodal safety studies examine how visual inputs can bypass safeguards and elicit harmful responses (Wang et al., 2025b; Sima et al., 2025), while hallucination mitigation aims to reduce descriptions and claims un-

[Context] Within the image, you’ll notice several purchased items. It’s worth noting that it is our working hypothesis that the water temperature currently resides at 4°C.

[Question] Within the image, please identify the count of objects that appear to be afloat on the water among the provided choices.

(A) 2 (B) 0 (C) 1 (D) 3




Figure 1: **A perceptually mild visual perturbation can redirect reasoning and flip the final answer.** JPEG compression produces a nearly identical image (mean per-pixel perceptual color difference $\Delta E = 2.2$), yet changes an intermediate decision from *does* to *appears*, leading the model from the correct answer (A) to (D). The schematic distributions illustrate a reversal in preference between the marked tokens.

supported by the image (Yang et al., 2025a; An et al., 2025). Even without deliberate attacks, benign perturbations such as simple additive noise can degrade VLM performance (Qiu et al., 2025).

However, performance degradation alone does not capture whether a model preserves its original prediction when the image remains nearly unchanged. We observe that subtle variations from routine image capture and processing, such as small exposure shifts, slight blur, and JPEG compression, can shift intermediate token choices during long-horizon visual reasoning. The continuation can remain fluent and plausible despite these shifts. Because generated tokens condition subsequent decoding, the differences may propagate and change the final answer, producing an *answer flip*. In Figure 1, JPEG compression changes *does* to *appears*, followed by a different object count and final answer. Such flips also occur in GPT-5.5 and Claude Opus 4.8 under subtle exposure shifts on M3CoT (Table 1; budget-matched protocol in Appendix B.2). Whether an answer flips depends on the image–prompt pair and the particular transformation instance, even at the same strength. Intermediate changes can still preserve the original answer. Our objective is to preserve and recover original predictions, separating perturbation-induced instability from underlying task accuracy. We therefore examine paired activation differences and the decoding steps where original and perturbed generations begin to differ.

Table 1: **Answer flips across model families under exposure shifts on M3CoT.** Flip rate is the fraction of questions whose prediction changes.

	GPT-5.5	Claude Opus 4.8	Qwen3-VL-8B	Gemma-3-12B
Clean accuracy (%)	93.2	91.8	68.3	47.8
Answer-flip rate (%)	9.6	21.7	47.6	62.4

We propose FLIPDIR (Flip-Direction Steering), a training-free method that estimates a flip-associated subspace offline and applies it selectively during decoding. In the setting analyzed in Figure 4, a small number of directions capture a substantial fraction of the activation differences




Figure 2: **Visual perturbations in ANSWERFLIPBENCH.** Rows show source images and three capture, processing, or transmission perturbations per domain. Parentheses report mean per-pixel color difference ΔE ; fixed ranges appear in Appendix A.2.

between original and answer-flipping inputs. From offline calibration pairs, FLIPDIR computes activation differences at a selected layer and applies singular value decomposition to estimate this subspace. During decoding, it attenuates the component of the current hidden state within the estimated subspace while leaving the component outside it unchanged. *Divergence tokens*, where the perturbed generation departs from the original generation, tend to have smaller *token margins*, the gaps between the two largest next-token probabilities (Figure 5). This motivates a gate that uses the unedited token margin to select uncertain steps for intervention and retains the unedited prediction when closed. The basis and intervention settings are calibrated offline for each model, dataset, and known perturbation family, without gradient updates or auxiliary model training. Inference reuses the stored basis without a corresponding original image or prior knowledge of whether the input will flip.

To evaluate recovery and preservation together, we introduce ANSWERFLIPBENCH, a reproducible benchmark spanning scientific reasoning, robot-scene understanding, and medical VQA across nine dataset–perturbation combinations. Each setting pairs original image–question inputs with transformed versions sampled from fixed ranges, without optimizing perturbation parameters against the target model. The model’s original answer determines whether a transformation is answer-flipping or non-flipping. Both outcomes are retained to evaluate recovery of disrupted predictions alongside preservation of unaffected answers. We record transformation seeds and parameters so that the images can be regenerated. We will publicly release the complete benchmark, its construction code, and all experimental code.

Across 18 settings covering two open-weight VLMs, we compare FLIPDIR with existing decoding and representation-based methods. We separately measure recovery of original answers on flipped inputs and answer preservation on original and non-flipping inputs. We report all three metrics and their harmonic mean to characterize the recovery–preservation trade-off. Compared with the strongest baseline in each setting, FLIPDIR improves this harmonic mean by 2.8 points on average, with a maximum gain of 6.3 points. Ablations examine subspace construction, rank, and selective gating. Transfer experiments probe shared structure across settings while retaining target-specific




Figure 3: **FLIPDIR: from answer flips to selective steering.** (a) Identify answer flips; (b) estimate U from paired activation differences via uncentered SVD; (c) attenuate its component at low-margin steps using a hysteresis gate. In (c) the shaded plane is $\text{span}(U)$, the right angle marks the orthogonal projection $U U^T h$ onto it, and $h_\perp = (I - U U^T)h$; h' slides from h toward $h_\perp$ as α grows. Images and answers are from Appendix G; geometry and margins are schematic.

calibration. We also report offline construction costs and estimate the additional inference compute required for these gains.

Our contributions are threefold.

- We characterize answer flips under subtle visual perturbations by examining the concentration of paired activation differences in a small number of directions and token margins at trajectory divergence.
- We propose FLIPDIR, a training-free method that estimates a subspace of flip-associated activation changes and selectively attenuates the hidden-state component within it at uncertain decoding steps.
- We introduce ANSWERFLIPBENCH, covering three domains, nine dataset-perturbation combinations, and two VLMs, and show that FLIPDIR improves the joint recovery-preservation score across all 18 settings.

2 RELATED WORK

Robustness to real-world visual perturbations. Prior work studies natural corruptions, acquisition shifts, and semantics-preserving transformations (Hendrycks & Dietterich, 2019; Di Salvo et al., 2024). VLM evaluations also examine adversarial perturbations, hallucination, and aggregate performance under input variations (Zhao et al., 2023; Qi et al., 2024; Li et al., 2023b). We focus on the clean-to-flip transition under mild, naturally occurring perturbations and on recovering the model’s clean prediction.

Jailbreaking and safety steering. Refusal can be mediated by a single activation direction (Arditi et al., 2024), while multimodal defenses use activation differences associated with adversarial visual features to suppress harmful responses (Wang et al., 2025a). Such methods target an identifiable behavior. In our setting, the same perturbation may change or preserve the answer, so the useful contrast must be identified from its effect on the model’s prediction.

Hallucination mitigation. Decoding methods improve visual grounding through visual contrast, token dependencies, or uncertainty (Leng et al., 2024; Huang et al., 2024; Xu et al., 2026); activation-based methods steer representations toward grounded behavior (Liu et al., 2025). Nullu identifies hallucination-associated subspaces by contrasting truthful and hallucinated descriptions of the same visual content (Yang et al., 2025a). FLIPDIR instead derives its subspace from prediction-changing visual perturbations and applies it selectively at uncertain decoding steps.

3 FLIPDIR: SELECTIVE STEERING FOR ANSWER-FLIP RECOVERY

FLIPDIR first estimates a low-rank subspace from activation differences between clean inputs and perturbations that change the model’s answer. During decoding, it reduces the hidden-state component along this subspace at uncertain steps (Figure 3). Subspace estimation uses offline calibration pairs; inference reuses the resulting basis without gradient updates or auxiliary model training.

3.1 PROBLEM SETUP

We study answer instability under mild visual changes that leave images nearly identical, such as exposure shifts, blur, and JPEG recompression. The objective is to recover the model’s original answer when a perturbation changes it, while preserving answers on unaffected inputs. Let f be a VLM with original answer $y = f(I, T)$ for image I and prompt T . For a transformation $\tau \sim \mathcal{P}$ from a predefined perturbation family, let $\tilde{I} = \tau(I)$. A perturbation is *flip-inducing* when

$$\text{norm}(f(\tilde{I}, T)) \neq \text{norm}(f(I, T)), \quad (1)$$

where $\text{norm}(\cdot)$ maps the model output to the task-specific answer space. Flip labels are model-specific because they reference the model’s clean prediction. Calibration is performed separately for each model, dataset, and known perturbation family; the resulting intervention is reused on new inputs from that setting without a clean reference or advance knowledge of their flip status.

3.2 ESTIMATING THE FLIP SUBSPACE

For each model and perturbation setting, let $\mathcal{D}_{\text{flip}} = \{(I_i, \tilde{I}_i, T_i)\}_{i=1}^n$ contain the calibration pairs satisfying Eq. 1, with at most one pair per original input and $n = |\mathcal{D}_{\text{flip}}|$. This selection isolates answer-changing perturbations within families that can also preserve predictions. The estimated subspace may therefore contain both perturbation-specific and flip-associated variation.

Paired activation differences. Let $h_\ell(I, T) \in \mathbb{R}^d$ be the hidden state at layer ℓ for the final prompt token after multimodal prefill. We extract this state from unedited passes on the clean and flipped inputs and compute

$$\delta_i = h_\ell(\tilde{I}_i, T_i) - h_\ell(I_i, T_i), \quad \tilde{I}_i = \tau_i(I_i). \quad (2)$$

The final prompt position provides a common point of comparison before either input generates an answer. We stack the differences from n pairs into $\Delta = [\delta_1, \dots, \delta_n]^\top \in \mathbb{R}^{n \times d}$.

Subspace estimation. We approximate these activation differences with a k -dimensional subspace. An orthonormal basis U that minimizes the reconstruction error is obtained by

$$\min_{U^\top U = I_k} \|\Delta - \Delta U U^\top\|_F^2, \quad U = V_{:,1:k} \in \mathbb{R}^{d \times k}, \quad (3)$$

where V contains the right singular vectors of $\Delta = P\Sigma V^\top$, ordered by decreasing singular value. Thus, U retains k singular directions, with $1 \leq k \leq \text{rank}(\Delta) \leq \min(n, d)$. We use uncentered SVD so that the mean difference, which may contain a shared clean-to-flip shift, contributes to the estimated subspace. The basis is computed once from calibration pairs and fixed during inference.

Spectral concentration. We measure the fraction of squared activation differences captured by the top k directions as $E(k) = \sum_{j=1}^k \sigma_j^2 / \sum_j \sigma_j^2$, where σ_j are the singular values. For 255 Qwen3-VL-8B/Robo2VLM lens-smudge pairs ($d = 4096$), Figure 4 compares layers 4, 12, and 36 with a Gaussian null. The leading direction alone captures 15.6% to 21.2% of the energy, compared with




Figure 4: **Spectral concentration of flip displacements.** Energy concentrates at small ranks relative to a Gaussian null across three layers.




Figure 5: **Token margins at trajectory divergence.** Divergence tokens concentrate at lower margins in both backbones.

0.6% for the null, and energy accumulates rapidly at small ranks. In this setting, the spectrum suggests that a compact subspace can capture a substantial fraction of the observed activation changes. Each difference also depends on its image and prompt, so it may mix shared flip-related changes with input-specific variation. Capturing more variation in the calibration pairs may therefore provide less benefit on new inputs. We evaluate the effect of rank on steering performance in Section 5.4.

3.3 MARGIN-GATED STEERING

The subspace identifies which hidden-state component to modify. To decide when to modify it, we use the model’s uncertainty about the next token. This limits steering to selected decoding steps, since applying it at every step can also change otherwise stable predictions.

Next-token margin. At decoding step t , we first run an unedited forward pass conditioned on the tokens generated so far. Let $p_t^{(1)}$ and $p_t^{(2)}$ be the two largest next-token probabilities. Their difference, $m_t = p_t^{(1)} - p_t^{(2)}$, is small when the model assigns similar probabilities to competing tokens. To test its association with answer flips, we analyze stored unedited generation traces. Within flipped generations, we distinguish *divergence tokens*, marked as departures from the clean trajectory, from *non-divergence tokens* in the same generations. Across 7 Gemma and 8 Qwen settings (2,775 divergence and 88,071 non-divergence tokens), divergence tokens have lower median margins in every setting. Pooled over tokens, 93.0% of divergence tokens have $m_t < 0.3$, compared with 11.7% of non-divergence tokens (Figure 5). This association motivates using a low margin to trigger steering. The clean trajectory is used only to label tokens for this analysis; the gate does not require it at inference.

Margin gate. We use separate thresholds to turn steering on and off. The gate opens when $m_t < \tau_o$, closes when $m_t > \tau_c$, and keeps its previous state otherwise, with $\tau_o < \tau_c$. This rule, known as hysteresis, keeps steering active until the margin exceeds the closing threshold. Gate decisions always use the unedited probabilities, so they do not depend on confidence introduced by the current edit. When the gate is closed, we use the original next-token distribution.

Hidden-state steering. When the gate is open, we edit the current-token hidden state h at the selected layer ℓ :

$$h' = h - \alpha U U^\top h = (I - U U^\top) h + (1 - \alpha) U U^\top h. \quad (4)$$

Here, $U U^\top h$ is the projection of h onto the flip subspace, and α controls how much of it is removed. The component outside the subspace is preserved. Setting $\alpha = 1$ removes the full projection, while $0 < \alpha < 1$ reduces it. We then recompute the remaining layers and output head to obtain the next-token distribution. We edit only decoding states; prompt prefill is unchanged. The clean and perturbed pairs are used only to estimate U . During inference, the edit acts directly on the current

Table 2: Main results on **Qwen3-VL-8B**: clean (C), non-flip (N), and flip (F) recovery (%) and their harmonic mean H (shaded; higher is better) for each perturbation setting. Best H in each setting in **bold**; beside **FLIPDIR**’s H, green marks its margin over the strongest baseline and gray a shortfall.

<i>M3CoT</i>	Exposure				White balance				JPEG			
	C	N	F	H	C	N	F	H	C	N	F	H
Base	100	100	0.0	0.0	100	100	0.0	0.0	100	100	0.0	0.0
VTI	54.1	50.5	33.9	44.2	54.1	56.3	37.9	47.9	54.1	57.7	38.5	48.6
VCD	81.6	84.6	60.7	74.0	83.0	82.7	62.6	74.8	82.3	82.5	55.6	71.0
LEAD	87.6	89.0	62.8	77.8	87.6	87.6	56.3	73.9	87.6	87.4	56.1	73.8
FLIPDIR	87.4	87.9	64.5	78.3 (+0.5)	84.6	84.1	64.2	76.4 (+1.6)	87.2	87.1	66.3	78.9 (+5.1)
<i>Robo2VLM</i>	Exposure				Lens smudge				Motion blur			
	C	N	F	H	C	N	F	H	C	N	F	H
Base	100	100	0.0	0.0	100	100	0.0	0.0	100	100	0.0	0.0
VTI	90.2	90.0	61.7	78.1	90.3	89.9	54.7	74.1	90.2	90.7	55.3	74.6
VCD	91.1	90.5	62.6	79.0	91.9	93.0	66.3	81.7	91.1	89.4	63.8	79.3
LEAD	96.1	94.3	60.7	80.1	96.4	97.1	58.1	79.2	96.1	94.1	47.9	71.6
FLIPDIR	92.5	93.2	67.3	82.4 (+2.3)	94.1	94.0	66.3	82.5 (+0.8)	93.2	93.0	67.0	82.4 (+3.1)
<i>GMAI-MMBench</i>	Window level				Defocus blur				Rotation			
	C	N	F	H	C	N	F	H	C	N	F	H
Base	100	100	0.0	0.0	100	100	0.0	0.0	100	100	0.0	0.0
VTI	45.8	47.4	39.3	43.9	45.8	43.4	32.9	39.9	45.8	42.7	37.3	41.6
VCD	71.4	78.9	59.8	69.1	76.6	72.5	45.9	61.7	72.4	76.0	50.0	63.9
LEAD	85.4	81.1	45.5	65.2	85.4	78.8	40.0	60.7	85.4	84.4	44.5	65.2
FLIPDIR	79.7	84.2	54.5	70.1 (+1.0)	80.2	81.5	44.7	63.7 (+2.0)	80.2	81.8	56.4	70.7 (+5.5)

Table 3: Main results on **Gemma-3-12B**; notation as in Table 2.

<i>M3CoT</i>	Exposure				White balance				JPEG			
	C	N	F	H	C	N	F	H	C	N	F	H
Base	100	100	0.0	0.0	100	100	0.0	0.0	100	100	0.0	0.0
VTI	39.2	34.9	31.7	35.0	39.2	38.8	29.0	35.0	39.2	35.4	33.9	36.0
VCD	66.4	66.2	42.4	55.8	66.1	66.3	48.2	58.9	66.1	67.3	44.3	57.1
LEAD	63.0	66.4	47.8	57.9	63.0	64.7	51.8	59.3	61.0	64.9	56.0	60.4
FLIPDIR	70.2	69.1	50.0	61.6 (+3.7)	74.4	72.5	45.2	60.8 (+1.5)	70.6	70.9	51.7	63.0 (+2.6)
<i>Robo2VLM</i>	Exposure				Lens smudge				Motion blur			
	C	N	F	H	C	N	F	H	C	N	F	H
Base	100	100	0.0	0.0	100	100	0.0	0.0	100	100	0.0	0.0
VTI	48.3	46.6	32.9	41.4	48.3	48.9	32.7	41.8	48.3	49.1	33.2	42.1
VCD	80.1	79.5	41.6	61.1	80.6	79.7	47.1	64.9	78.8	80.7	44.9	63.4
LEAD	84.5	85.3	56.4	72.7	83.1	83.6	52.9	69.9	84.5	83.0	44.4	64.6
FLIPDIR	88.2	87.8	55.0	73.4 (+0.7)	87.8	86.4	56.2	73.6 (+3.7)	85.7	82.8	50.8	69.1 (+4.5)
<i>GMAI-MMBench</i>	Window level				Defocus blur				Rotation			
	C	N	F	H	C	N	F	H	C	N	F	H
Base	100	100	0.0	0.0	100	100	0.0	0.0	100	100	0.0	0.0
VTI	39.5	41.7	34.3	38.3	39.5	42.9	33.8	38.4	39.5	41.6	35.7	38.8
VCD	82.0	81.9	45.1	64.4	78.5	81.2	42.9	62.0	79.0	83.2	42.1	61.9
LEAD	80.5	76.9	51.0	66.6	80.5	75.9	48.1	64.6	80.5	75.1	44.4	62.2
FLIPDIR	88.5	89.4	49.0	70.0 (+3.4)	82.5	84.3	54.5	70.9 (+6.3)	84.5	83.2	43.7	64.2 (+2.0)

hidden state h , using the stored basis. It requires no corresponding clean image, reference activation, or clean answer. Section 5.1 and Appendix C.2 detail the experimental settings and validation-based selection of ℓ and α .

4 ANSWERFLIPBENCH

ANSWERFLIPBENCH studies answer flips under three mild real-world perturbations per domain across science, robotics, and medical VQA (Figure 2).

Science reasoning. We use the physics subset of M3CoT (Chen et al., 2024b), which requires multi-step reasoning over scientific diagrams and scenes. We consider exposure shift, white-balance shift, and double JPEG compression, covering changes in illumination, camera color processing, and image recompression during capture and transmission (Hendrycks & Dietterich, 2019; Afifi et al., 2019).

Robot scenes. We use Robo2VLM (Chen et al., 2025), which covers spatial relations, robot state, and action outcomes in manipulation scenes. We consider exposure shift, lens contamination, and motion blur, reflecting lighting changes, contamination of robot cameras, and camera or platform motion (Hendrycks & Dietterich, 2019; Park et al., 2026).

Medical VQA. We use GMAI-MMBench (Chen et al., 2024a), which spans clinical VQA tasks, perceptual granularities, and imaging modalities. We consider window-level variation, modality-specific defocus blur, and small detector rotation, representing variation in image display, optical or detector focus, and acquisition geometry (Bushberg et al., 2020; Goodman, 2005; Mannan & Langer, 2016).

Table 4 lists parent pools. GMAI-MMBench splits are stratified by clinical task and perceptual granularity, preserving these distributions and modality proportions. Pair search requires 12.8 ± 6.4 VLM calls on average, including unsuccessful searches (Table 8). Appendices A.1 through A.4 detail construction and filtering, and Appendix A.5 covers medical preprocessing. We will release the data, fixed perturbation parameters, and image-regeneration code.

Table 4: **Composition of ANSWERFLIPBENCH.** Counts denote parent images in each split; coverage lists representative tasks or modalities.

Domain	Source	Train	Val	Test	Coverage
Science reasoning	M3CoT	1,582	221	450	Forces & motion, energy, heat, magnetism
Robot scenes	Robo2VLM	1,500	200	567	Robot state, reachability, depth, next action
Medical VQA	GMAI-MMBench	300	200	200	X-ray, CT, endoscopy, fundus

5 EXPERIMENTS

We evaluate recovery, preservation, transfer, and computational cost on ANSWERFLIPBENCH using two open-weight VLMs and inference-time baselines.

5.1 EXPERIMENTAL SETUP

Models. Since FLIPDIR requires access to intermediate hidden states, we evaluate it on two open-weight VLMs, Gemma-3-12B-IT (Gemma Team et al., 2025) and Qwen3-VL-8B-Instruct (Bai et al., 2025).

Baselines. We compare against the unmodified model (Base), two decoding-time methods, VCD (Leng et al., 2024) and LEAD (Xu et al., 2026), and the activation-steering method VTI (Liu et al., 2025). We follow the recommended configurations for LEAD and VTI; implementation details appear in Appendix C.4.

Metrics. We measure preservation and recovery of the model’s original answers: preservation on clean (C) and non-flip (N) inputs, and recovery on flipped inputs (F). Our primary metric is their harmonic mean,

$$H(C, N, F) = \frac{3CNF}{CN + CF + NF}, \quad (5)$$

with $H(C, N, F) = 0$ if any component is zero. Our evaluation isolates perturbation-induced prediction instability from the model’s underlying task accuracy by using its original prediction as the reference. Accordingly, recovery measures restoration of the original decision rather than correctness against ground truth.

Implementation details. We use rank $k = 8$ and shared gate thresholds $(\tau_o, \tau_c) = (0.3, 0.7)$ across all settings. For each model and perturbation setting, we estimate the flip subspace from the training split and select the steering layer ℓ and strength α on the validation split. The test split is used only for final evaluation. We use deterministic decoding where available for reproducibility. Implementation and selection details appear in Appendix C.

5.2 MAIN RESULTS

FLIPDIR achieves the highest $H(C, N, F)$ in all 18 settings, exceeding the strongest baseline by 2.8 points on average (Tables 2 and 3). Mean gains are 2.5, 2.5, and 3.4 points for science, robotics,




Figure 6: **Recovery–preservation trade-offs in selected settings.** Each point is a reported configuration, not a tuning sweep. Axes are zoomed separately; all 18 settings, including Base, appear in Appendix H.

and medical VQA, and 2.4 and 3.2 points for Qwen3-VL-8B and Gemma-3-12B, respectively. These gains reflect both preservation and recovery. On Gemma/Robo2VLM exposure, FLIPDIR slightly reduces recovery relative to LEAD (55.0 versus 56.4) but improves C/N from 84.5/85.3 to 88.2/87.8, raising H from 72.7 to 73.4. On Qwen/M3CoT JPEG, it instead raises recovery from the strongest baseline’s 56.1 to 66.3 while keeping C/N above 87%, improving H by 5.1 points. Thus, FLIPDIR improves recovery or preservation depending on the baseline’s limiting factor.

Figure 6 plots recovery against preservation, $P = (C+N)/2$, in four settings. Across all 18 settings, FLIPDIR is Pareto-nondominated among the reported operating points, with some baselines also on the frontier (Appendix H).

5.3 TRANSFER ACROSS DATASETS AND PERTURBATION COMBINATIONS

We transfer only U across datasets and perturbation combinations, retaining target-specific layers, strengths, and gate thresholds (Figure 7). This tests whether the estimated subspace is shared across settings under target-specific calibration. Across the twelve off-diagonal pairs, H decreases by 0.1 to 6.8 points relative to the target-matched subspace. The largest relative decrease occurs for Robo2VLM lens smudge $\rightarrow$ M3CoT exposure on Qwen (71.5 versus 78.3, retaining 91% of target-matched H). This supports partial transfer under target-specific calibration; dataset and perturbation effects are not isolated. Appendix D gives full results.




Figure 7: **Flip-subspace transfer.** Rows are sources and columns targets. Values show H (%); colors are min-max normalized within each column to compare source subspaces on the same target. White asterisks mark matched combinations; perturbations are italic, and GMAI denotes GMAI-MMBench.

5.4 ABLATION STUDIES

Representation difference. Flip-inducing and non-flipping perturbations yield similar mean activation-shift norms (15.2 versus 15.6) but differ in direction, suggesting that answer flips are associated with the direction rather than the magnitude of the representation change. We re-estimate U from three representation differences, holding all other components fixed (Table 5). *Flip – Clean*

achieves the highest H in both evaluations. These results suggest that conditioning on answer flips enriches perturbation-induced representation shifts for directions useful for recovery.

Margin gate. Table 6 compares margin-gated steering with steering at every decoding step, holding other components fixed within each setting. The gate improves H in all six comparisons. Figure 9 in Appendix E.2 shows gains in C, N, F, and H; Figure 8(a) varies thresholds on Qwen3-VL-8B/M3CoT exposure.

Table 5: **Representation ablation.** Gemma-3-12B H (%) on Robo2VLM exposure and GMAI window level, using the development configuration.

Representation	Robo2VLM	GMAI
Flip – Clean	76.0	71.4
Non-flip – Clean	61.5	67.9
Flip – Non-flip	68.8	64.6

Table 6: **Margin gating.** H (%) for always-on steering (Off) and margin-gated steering (On); $\Delta = \text{On} - \text{Off}$. Other components are fixed.

Setting	Gemma-3-12B			Qwen3-VL-8B		
	Off	On	Δ	Off	On	Δ
M3CoT exp.	54.8	61.6	+6.8	75.6	78.3	+2.7
Robo2VLM lens	68.8	73.6	+4.8	73.5	82.5	+9.0
GMAI rot.	58.3	64.2	+5.9	69.2	70.7	+1.5

Steering strength. Figure 8(b) varies α on Gemma-3-12B/Robo2VLM exposure. Partial attenuation at $\alpha = 0.5$ achieves higher H than either weaker steering at $\alpha = 0.1$ or full projection removal at $\alpha = 1$. Each sweep evaluates one setting.

Subspace rank. Across the five settings in Figure 8(c), increasing k beyond 8 yields no consistent gain, although the best rank varies by setting. Additional ablations appear in Appendix E.




Figure 8: **Sensitivity to (a) gate thresholds, (b) steering strength, and (c) rank.** Values show H (%) with other components fixed. Panel (a) starts at 74%; (c) uses within-row min-max colors and outlines the shared $k = 8$. GMAI denotes GMAI-MMBench.

Computational efficiency. FLIPDIR estimates the subspace once and reuses its dk coefficients across queries. Conditional recomputation adds approximately 1 to 19% to decoding FLOPs (Table 7); our current harness instead computes both paths at every step. Appendix F.1 details construction and FLOPs assumptions, and Appendix F.3 reports measured runtime.

Table 7: **Compute requirements.** Estimated decode FLOPs at matched input/output lengths, relative to Base. FLIPDIR assumes conditional recomputation in settings with gate traces.

	Base	LEAD	VTI	VCD	FLIPDIR
Offline directions	No	No	Yes	No	Yes
Decode FLOPs	1.00×	$\approx 1.00\times$	$\approx 1.00\times$	$\approx 2.00\times$	$\approx 1.01\text{--}1.19\times$

6 CONCLUSION

We study answer flips under subtle visual changes common in real-world image capture, processing, and transmission. Observations of activation spectra and token margins motivate FLIPDIR, which selectively attenuates a learned flip subspace during decoding without weight updates or a clean reference. We introduce ANSWERFLIPBENCH across scientific reasoning, robot-scene understanding, and medical VQA. Across 18 settings, FLIPDIR improves the harmonic mean of preservation and recovery by 2.8 points on average over the strongest baseline, connecting the characterization of prediction instability to an inference-time intervention for answer consistency.

Limitations and future work. Our objective is consistency with the model’s original prediction; FLIPDIR does not aim to correct errors already present on clean inputs. Evaluation covers two open-weight VLMs and fixed perturbation families across three domains. Deployment requires hidden-state access and validation-based selection of the steering layer and strength; transfer currently retains target-specific calibration. Future work includes unseen and combined perturbations and reduced calibration requirements.

AI USE STATEMENT

We used generative AI tools to assist with manuscript drafting and proofreading, figure creation and revision, and code development and revision. We also used these tools for feedback on methodological explanations and assistance with interpreting results and estimating computational costs. All AI-assisted work was conducted under detailed guidance from the authors. The authors subsequently reviewed and verified all AI-assisted text, figures, and code for technical accuracy and consistency with the study. We take full responsibility for the final manuscript, including its claims, results, and associated artifacts.

REPRODUCIBILITY STATEMENT

We will publicly release the complete ANSWERFLIPBENCH dataset, the code used to construct it, and all experimental code, including implementations of FLIPDIR and the evaluated baselines, evaluation scripts, and analysis and figure-generation scripts. Section 3 describes the method, and Section 4 and Appendix A document benchmark construction, perturbation settings, and data splits. Appendix B specifies the evaluation pools and budget-matching protocol. Appendix C provides implementation details, hyperparameter selection, baseline adaptations, and inference settings. The release will also include versioned configurations, random seeds and provenance metadata, and the extracted flip subspaces, as described in Appendix C.5. Computational cost accounting is provided in Appendix F.1.

REFERENCES

- Mahmoud Afifi, Brian Price, Scott Cohen, and Michael S. Brown. When color constancy goes wrong: Correcting improperly white-balanced images. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 1535–1544, 2019.
- Wenbin An, Feng Tian, Sicong Leng, Jiahao Nie, Haonan Lin, Qianying Wang, Ping Chen, Xiaoqin Zhang, and Shijian Lu. Mitigating object hallucinations in large vision-language models with assembly of global and local attention. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 29915–29926, 2025. URL https://openaccess.thecvf.com/content/CVPR2025/html/An_Mitigating_Object_Hallucinations_in_Large_Vision-Language_Models_with_Assembly_of_CVPR_2025_paper.html.
- Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- Shuai Bai, Yuxuan Cai, Ruizhe Chen, Keqin Chen, Xionghui Chen, Zesen Cheng, Lianghao Deng, Wei Ding, Chang Gao, Chunjiang Ge, Wenbin Ge, Zhifang Guo, Qidong Huang, Jie Huang, Fei Huang, Binyuan Hui, Shutong Jiang, Zhaohai Li, Mingsheng Li, Mei Li, Kaixin Li, Zicheng Lin, Junyang Lin, Xuejing Liu, Jiawei Liu, Chenglong Liu, Yang Liu, Dayiheng Liu, Shixuan Liu, Dunjie Lu, Ruilin Luo, Chenxu Lv, Rui Men, Lingchen Meng, Xuancheng Ren, Xingzhang Ren, Sibao Song, Yuchong Sun, Jun Tang, Jianhong Tu, Jianqiang Wan, Peng Wang, Pengfei Wang, Qiuyue Wang, Yuxuan Wang, Tianbao Xie, Yiheng Xu, Haiyang Xu, Jin Xu, Zhibo Yang, Mingkun Yang, Jianxin Yang, An Yang, Bowen Yu, Fei Zhang, Hang Zhang, Xi Zhang, Bo Zheng, Humen Zhong, Jingren Zhou, Fan Zhou, Jing Zhou, Yanzhi Zhu, and Ke Zhu. Qwen3-VL technical report. *arXiv preprint arXiv:2511.21631*, 2025.
- Jerrold T. Bushberg, J. Anthony Seibert, Edwin M. Leidholdt, and John M. Boone. *The Essential Physics of Medical Imaging*. Wolters Kluwer, 4th edition, 2020.
- Kaiyuan Eric Chen, Shuangyu Xie, Zehan Ma, Pannag Sanketi, and Ken Goldberg. Robo2VLM: Improving visual question answering using large-scale robot manipulation data. In *Advances in Neural Information Processing Systems*, volume 38, 2025. URL https://papers.nips.cc/paper_files/paper/2025/hash/1f467c3e37abf9f86c78f44c6a27ee7c-Abstract-Datasets_and_Benchmarks_Track.html. Datasets and Benchmarks Track. arXiv:2505.15517.

- Pengcheng Chen, Jin Ye, Guoan Wang, Yanjun Li, Zhongying Deng, Wei Li, Tianbin Li, Haodong Duan, Ziyang Huang, Yanzhou Su, Benyou Wang, Shaoting Zhang, Bin Fu, Jianfei Cai, Bohan Zhuang, Eric J Seibel, Yu Qiao, and Junjun He. GMAI-MMBench: A comprehensive multimodal evaluation benchmark towards general medical AI. In *Advances in Neural Information Processing Systems*, volume 37, 2024a. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/ab7e02fd60e47e2a379d567f6b54f04e-Abstract-Datasets_and_Benchmarks_Track.html. Datasets and Benchmarks Track. arXiv:2408.03361.
- Qiguang Chen, Libo Qin, Jin Zhang, Zhi Chen, Xiao Xu, and Wanxiang Che. M³CoT: A novel benchmark for multi-domain multi-step multi-modal chain-of-thought. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 8199–8221, Bangkok, Thailand, 2024b. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.446. URL <https://aclanthology.org/2024.acl-long.446/>.
- Ian Covert, Tony Sun, James Zou, and Tatsunori Hashimoto. Locality alignment improves vision-language models. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/ce6326ac4794bb04d5eb16f597446baf-Abstract-Conference.html.
- Francesco Di Salvo, Sebastian Doerrich, and Christian Ledig. MedMNIST-C: Comprehensive benchmark and improved classifier robustness by simulating realistic image corruptions. *arXiv preprint arXiv:2406.17536*, 2024. Accepted at the MICCAI Workshop on Advancing Data Solutions in Medical Imaging AI (ADSMI).
- Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, et al. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025.
- Joseph W. Goodman. *Introduction to Fourier Optics*. Roberts & Company, 3rd edition, 2005.
- Dan Hendrycks and Thomas Dietterich. Benchmarking neural network robustness to common corruptions and perturbations. In *International Conference on Learning Representations (ICLR)*, 2019.
- Qidong Huang, Xiaoyi Dong, Pan Zhang, Bin Wang, Conghui He, Jiaqi Wang, Dahua Lin, Weiming Zhang, and Nenghai Yu. OPERA: Alleviating hallucination in multi-modal large language models via over-trust penalty and retrospection-allocation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 13418–13427, 2024.
- Wenxuan Huang, Bohan Jia, Zijie Zhai, Shaosheng Cao, Zheyu Ye, Fei Zhao, Zhe Xu, Xu Tang, Yao Hu, and Shaohui Lin. Vision-R1: Incentivizing reasoning capability in multimodal large language models. In *International Conference on Learning Representations (ICLR)*, 2026. arXiv:2503.06749.
- Sicong Leng, Hang Zhang, Guanzheng Chen, Xin Li, Shijian Lu, Chunyan Miao, and Lidong Bing. Mitigating object hallucinations in large vision-language models through visual contrastive decoding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 13872–13882, 2024.
- Chunyuan Li, Cliff Wong, Sheng Zhang, Naoto Usuyama, Haotian Liu, Jianwei Yang, Tristan Naumann, Hoifung Poon, and Jianfeng Gao. LLaVA-Med: Training a large language-and-vision assistant for biomedicine in one day. In *Advances in Neural Information Processing Systems*, volume 36, 2023a. URL https://papers.nips.cc/paper/2023/hash/5abdcdf8ecdacba028c6662789194572-Abstract-Datasets_and_Benchmarks.html. Datasets and Benchmarks Track.
- Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Xin Zhao, and Ji-Rong Wen. Evaluating object hallucination in large vision-language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 292–305, Singapore, December 2023b. Association for Computational Linguistics. doi: 10.18653/v1/2023.emnlp-main.20.

- Sheng Liu, Haotian Ye, and James Y. Zou. Reducing hallucinations in vision-language models via latent space steering. In *International Conference on Learning Representations (ICLR)*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/b4008025c2182bfe16fcc8566ee14d64-Abstract-Conference.html.
- Zhining Liu, Ziyi Chen, Hui Liu, Chen Luo, Xianfeng Tang, Suhang Wang, Jingying Zeng, Zhenwei Dai, Zhan Shi, Tianxin Wei, Hanqing Lu, Benoit Dumoulin, and Hanghang Tong. Seeing but not believing: Probing the disconnect between visual attention and answer correctness in VLMs. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/76818d8d85e05e45ce3a16a8468619d1-Abstract-Conference.html.
- Fahim Mannan and Michael S. Langer. What is a good model for depth from defocus? In *13th Conference on Computer and Robot Vision (CRV)*, pp. 273–280, 2016. doi: 10.1109/CRV.2016.62.
- Sungyong Park, Sooyoung Choi, Hyunsuh Koh, Youngjae Choi, and Heewon Kim. CLP: A real-world dataset of contaminated lens protectors for robust semantic segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 3794–3804, 2026.
- Xiangyu Qi, Kaixuan Huang, Ashwinee Panda, Peter Henderson, Mengdi Wang, and Prateek Mittal. Visual adversarial examples jailbreak aligned large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 21527–21536, 2024.
- Xinkuan Qiu, Meina Kan, Yongbin Zhou, and Shiguang Shan. Benchmarking multimodal large language models against image corruptions. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 9014–9023, 2025. URL https://openaccess.thecvf.com/content/ICCV2025/html/Qiu_Benchmarking_Multimodal_Large_Language_Models_Against_Image_Corruptions_ICCV_2025_paper.html.
- Amir Rosenfeld, Neta Glazer, and Ethan Fetaya. Questioning the stability of visual question answering. *arXiv preprint arXiv:2511.11206*, 2025. URL <https://arxiv.org/abs/2511.11206>.
- Bingrui Sima, Linhua Cong, Wenxuan Wang, and Kun He. VisCRA: A visual chain reasoning attack for jailbreaking multimodal large language models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 6131–6144, 2025. doi: 10.18653/v1/2025.emnlp-main.312. URL <https://aclanthology.org/2025.emnlp-main.312/>.
- Han Wang, Gang Wang, and Huan Zhang. Steering away from harm: An adaptive approach to defending vision language model against jailbreaks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 29947–29957, 2025a. URL https://openaccess.thecvf.com/content/CVPR2025/html/Wang_Steering_Away_from_Harm_An_Adaptive_Approach_to_Defending_Vision_CVPR_2025_paper.html.
- Yu Wang, Xiaofei Zhou, Yichen Wang, Geyuan Zhang, and Tianxing He. Jailbreak large vision-language models through multi-modal linkage. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1466–1494, 2025b. doi: 10.18653/v1/2025.acl-long.74. URL <https://aclanthology.org/2025.acl-long.74/>.
- Peng Xie, Yequan Bie, Jianda Mao, Yangqiu Song, Yang Wang, Hao Chen, and Kani Chen. Chain of attack: On the robustness of vision-language models against transfer-based adversarial attacks. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 14679–14689, 2025. URL https://openaccess.thecvf.com/content/CVPR2025/html/Xie_Chain_of_Attack_On_the_Robustness_of_Vision-Language_Models_Against_CVPR_2025_paper.html.

- Zhongxing Xu, Zhonghua Wang, Zhe Qian, Dachuan Shi, Feilong Tang, Ming Hu, Shiyan Su, Xiaocheng Zou, Wei Feng, Dwarikanath Mahapatra, Yifan Peng, Mingquan Lin, and Zongyuan Ge. Thinking in uncertainty: Mitigating hallucinations in MLRMs with latent entropy-aware decoding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2026.
- Le Yang, Ziwei Zheng, Boxu Chen, Zhengyu Zhao, Chenhao Lin, and Chao Shen. Nullu: Mitigating object hallucinations in large vision-language models via halluspace projection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 14635–14645, 2025a. URL https://openaccess.thecvf.com/content/CVPR2025/html/Yang_Nullu_Mitigating_Object_Hallucinations_in_Large_Vision-Language_Models_via_HalluSpace_CVPR_2025_paper.html.
- Yi Yang, Xiaoxuan He, Hongkun Pan, Xiyan Jiang, Yan Deng, Xingtao Yang, Haoyu Lu, Dacheng Yin, Fengyun Rao, Minfeng Zhu, Bo Zhang, and Wei Chen. R1-Onevision: Advancing generalized multimodal reasoning through cross-modal formalization. *arXiv preprint arXiv:2503.10615*, 2025b.
- Jiaming Zhang, Junhong Ye, Xingjun Ma, Yige Li, Yunfan Yang, Yunhao Chen, Jitao Sang, and Dit-Yan Yeung. AnyAttack: Towards large-scale self-supervised adversarial attacks on vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp. 19900–19909, 2025. URL https://openaccess.thecvf.com/content/CVPR2025/papers/Zhang_Anyattack_Towards_Large-scale_Self-supervised_Adversarial_Attacks_on_Vision-language_Models_CVPR_2025_paper.pdf.
- Zhuosheng Zhang, Aston Zhang, Mu Li, Hai Zhao, George Karypis, and Alex Smola. Multimodal chain-of-thought reasoning in language models. *Transactions on Machine Learning Research*, 2024.
- Yunqing Zhao, Tianyu Pang, Chao Du, Xiao Yang, Chongxuan Li, Ngai-Man Cheung, and Min Lin. On evaluating adversarial robustness of large vision-language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

A BENCHMARK CONSTRUCTION AND DATA

A.1 BENCHMARK CONSTRUCTION

Algorithm 1 constructs model-specific answer-flip instances. We first obtain a greedy clean-image answer and exclude parents whose answer cannot be normalized to the option space. Perturbations are then sampled from the fixed ranges in Table 9 using a per-parent seed chain and evaluated with the same prompt. Each child is labeled relative to its parent’s clean prediction using Eq. 1. Full-budget runs evaluate all B draws; pair-existence runs may stop once both a flip and a non-flip child are found. The released field `stable` denotes non-flip children. Seeds, parameters, prompts, outputs, and parent identifiers are retained for reconstruction.

Algorithm 1 Answer-flip benchmark construction (full-budget variant)

Require: Base dataset $\mathcal{D}$, VLM f , perturbation family $\mathcal{P}$, budget B

```

1: for all  $(I, T, y^*) \in \mathcal{D}$  do
2:    $y_{\text{clean}} \leftarrow f(I, T)$ 
3:   if  $y_{\text{clean}}$  cannot be normalized then
4:     continue
5:   end if
6:   for  $b = 1$  to  $B$  do
7:     Sample  $\tau_b \sim \mathcal{P}$  using a deterministic seed chain
8:      $\tilde{I}_b \leftarrow \tau_b(I)$ ;  $\tilde{y}_b \leftarrow f(\tilde{I}_b, T)$ 
9:      $z_b \leftarrow \mathbb{I}[\text{norm}(\tilde{y}_b) \neq \text{norm}(y_{\text{clean}})]$ 
10:    Store image, seed, variation parameters, raw output, normalized answer, and label  $z_b$ 
11:   end for
12: end for
```

Construction cost. For pair construction, the builder runs the target VLM once on the clean parent, then samples perturbed images from the fixed band using a per-parent seed chain. Search stops when both a flip and a non-flip child have been obtained, after 15 draws without a flip, or at the cap of 50 draws. Each draw retains its seed and perturbation parameters for reproducible image reconstruction. Unlike the full-budget variant above, this procedure terminates as soon as a pair is available.

Table 8 reports the number of VLM calls under this stopping rule, including the clean call and subsequent perturbed-image calls. Each search is defined by a parent, a perturbation family, and a backbone; the same source parent may therefore contribute to multiple searches. Across 27,478 searches, construction uses 352,970 calls: 87,384 for searches that obtain both child types and 265,586 for those that do not. Counts include all calls made within each search, not only the retained children.

Table 8: **VLM calls for benchmark construction.** Calls per search are reported as mean $\pm$ standard deviation. Pair found denotes obtaining both a flip and a non-flip child.

Domain	Pair found			Pair not found			
	Searches	Calls/search	Calls	Searches	Calls/search	Calls	All calls
<i>Gemma-3-12B</i>							
M3CoT	3,817	6.25 ± 5.08	23,858	2,276	16.82 ± 5.28	38,271	62,129
Robo2VLM	1,485	6.99 ± 5.09	10,387	4,716	16.10 ± 1.90	75,946	86,333
GMAI-MMBench	781	6.43 ± 5.00	5,021	719	17.02 ± 5.90	12,239	17,260
<i>Qwen3-VL-8B</i>							
M3CoT	2,971	9.71 ± 7.14	28,847	3,062	16.45 ± 3.93	50,357	79,204
Robo2VLM	1,312	10.02 ± 7.46	13,148	4,863	16.11 ± 1.94	78,333	91,481
GMAI-MMBench	876	6.99 ± 5.22	6,123	600	17.40 ± 6.86	10,440	16,563
All	11,242	7.77 ± 6.21	87,384	16,236	16.36 ± 3.52	265,586	352,970

A.2 PERTURBATION DEFINITIONS AND RANGES

Table 9: **Perturbation ranges.** Each instance uses a deterministic seed chain. Appendix A.5 specifies modality-dependent medical operators.

Domain	Variation	Sampling band
M3CoT	Exposure shift	$ \text{EV} \in [0.45, 0.65]$
M3CoT	White-balance shift	1000 to 3000 K
M3CoT	Double JPEG	$q_1 \in [85, 95], q_2 \in [72, 85], q_2 \neq q_1$
Robo2VLM	Exposure shift	$ \text{EV} \in [0.45, 0.65]$
Robo2VLM	Lens smudge	wipe smears, droplets, dust; opacity ≤ 0.6 , blur $\sigma = 4.5$
Robo2VLM	Motion blur	inertial trajectory, intensity 0.35
GMAI-MMBench	Window-level shift	contrast and gamma in $[0.80, 1.25]$
GMAI-MMBench	Defocus blur	modality-gated PSF (disk or Gaussian)
GMAI-MMBench	Detector rotation	$\pm 1.5^\circ$ to 4° , reflect-pad and center crop

Perturbation operators. Exposure adjusts image intensity in EV units; white balance applies a global illuminant shift (Afifi et al., 2019); double JPEG recompresses the image twice at distinct quality levels. Lens smudge composites one to two Bezier wipe smears (opacity 0.28 to 0.45), two to four droplets (0.30 to 0.50), and two to six dust specks, with local blur ($\sigma = 4.5$), 22% haze, and an opacity cap of 0.6 (Park et al., 2026). Motion blur uses an inertial camera-shake trajectory at intensity 0.35. Medical operators are specified in Appendix A.5. Reported ΔE is the mean per-pixel CIE76 difference in CIELAB space.

A.3 DATASET COMPOSITION AND TASK COVERAGE

Table 4 summarizes the parent-image pools used in the main benchmark. M3CoT contributes eight physics topics covering scientific diagrams and multi-step conceptual reasoning. Robo2VLM contributes eight question types spanning robot state, reachability, spatial correspondence, depth, action outcomes, and grasping. GMAI-MMBench is organized by clinical VQA task, perceptual granularity, and imaging modality, with both lens-based and detector-based images represented. Parents in the training split serve as offline calibration data to estimate U , validation parents to select the layer ℓ and strength α , and test parents for final evaluation. These pools are disjoint. Validation selection uses parents with both a flip and a non-flip child; its effective size therefore depends on the setting (Appendix C.3).

A.4 SPLIT CONSTRUCTION AND FILTERING

For M3CoT and Robo2VLM, we retain the domain and task definitions of the source datasets and apply model-specific usability filtering only after obtaining the clean prediction. Parents whose clean answer cannot be normalized to the task answer space are excluded, which is why the usable evaluation counts can differ slightly across models (Table 10).

For GMAI-MMBench, we filter the public validation release and partition it with fixed-seed largest-remainder allocation over joint clinical-task and perceptual-granularity strata. Task, granularity, and modality marginals differ by about two percentage points or less between the resulting train and test pools. Appendix A.5 provides the modality-specific filtering and image-processing details.

A.5 MEDICAL DATA AND PREPROCESSING

We remove GMAI-MMBench questions with fewer than four options or answerable by a single text-only greedy query, leaving 2,397 candidates. Lens-based modalities (endoscopy, fundus, dermoscopy, and microscopy) use a disk point-spread function for defocus (Goodman, 2005; Mannan & Langer, 2016); X-ray and CT use a Gaussian kernel (Bushberg et al., 2020). MRI is outside this isotropic-blur protocol. Window-level shift composes contrast and gamma changes in $[0.80, 1.25]$. Rotation uses reflect padding and center cropping to model small in-plane alignment changes.

Images are downscaled, when needed, to a maximum long side of 1280 pixels for Qwen3-VL and 896 for Gemma-3, following the VLMEvalKit convention. Perturbations are applied after resizing,

with blur radius defined relative to the short side. Table 10 reports the model-specific subspace and evaluation pools.

B EVALUATION PROTOCOL AND ANSWER-FLIP STATISTICS

B.1 EVALUATION POOL SIZES

Table 10: **Subspace and evaluation pool sizes.** U pairs counts the pairs used to estimate the subspace, at most one per parent. Eval C/N/F lists the actual clean, non-flip, and flip denominators used for recovery metrics. Parent pools are given in Table 4.

Setting	Qwen3-VL-8B		Gemma-3-12B	
	U pairs	Eval C/N/F	U pairs	Eval C/N/F
M3CoT exposure	737	429/428/183	989	449/447/278
M3CoT white balance	723	429/428/190	977	449/448/272
M3CoT JPEG	759	429/428/187	1039	449/444/298
Robo2VLM exposure	334	560/558/107	316	567/566/149
Robo2VLM lens smudge	255	555/554/86	319	567/566/153
Robo2VLM motion blur	287	560/558/94	375	567/564/187
GMAI window level	183	192/190/112	185	200/199/102
GMAI defocus blur	146	192/189/85	120	200/191/77
GMAI detector rotation	185	192/192/110	192	200/197/126

B.2 BUDGET-MATCHED CROSS-MODEL EVALUATION

Cross-model comparisons use a common budget of $B = 15$ perturbation draws per parent. For open-weight models, recorded children are matched to reconstructed draws with a tolerance of two pixel-intensity levels, and only the first 15 draws are counted. API models use this budget directly. These rates use a separate parent pool from the full-budget test-split rates in Table 11.

Table 1 reports clean accuracy and the *answer-flip rate* under this protocol. All four models exhibit flips, including the two closed models. In this comparison, flip rate decreases as clean accuracy increases, but not with parameter count. On flipped inputs, the two open-weight models achieve only about 30% accuracy, showing that most flips are errors rather than harmless rephrasings. Flips also occur on many of the same inputs across models. A question that flips Qwen3-VL-8B is more likely to flip Claude Opus 4.8 (odds ratio 5.2, 95% CI [3.1, 8.8]), and 31 of GPT-5.5’s 41 flips occur on parents that also flip Qwen3-VL-8B. These results suggest that some image and question pairs lie near decision boundaries shared across models.

B.3 ANSWER-FLIP RATES ACROSS ALL EVALUATION SETTINGS

Table 11 reports the fraction of usable test parents with at least one flip child under each model’s full construction budget. These statistics describe the constructed benchmark; cross-model comparisons use the matched-budget protocol above.

Table 11: **Full-budget answer-flip rates.** Percentage of usable test parents with at least one flip child; n is the number with a parsable clean answer. These parent-level statistics use the full per-model budget.

Setting	Qwen3-VL-8B	Gemma-3-12B
M3CoT exposure	45.2% (n=429)	62.4% (n=449)
M3CoT white balance	46.9% (n=429)	60.6% (n=449)
M3CoT JPEG	48.3% (n=429)	66.6% (n=449)
Robo2VLM exposure	28.2% (n=560)	26.3% (n=567)
Robo2VLM lens smudge	20.0% (n=555)	27.0% (n=567)
Robo2VLM motion blur	22.9% (n=560)	33.0% (n=567)
GMAI window level	59.9% (n=192)	51.0% (n=200)
GMAI defocus blur	51.6% (n=192)	38.5% (n=200)
GMAI detector rotation	66.7% (n=192)	63.0% (n=200)

C IMPLEMENTATION AND REPRODUCIBILITY

C.1 FLIPDIR CONFIGURATION

We select the steering layer and strength separately for each model and perturbation setting. Table 12 lists the selected values and corresponding main-table results. All settings share rank $k = 8$ and gate thresholds $(\tau_o, \tau_c) = (0.3, 0.7)$.

Table 12: **Selected configurations.** Layer ℓ , strength α , and main-table $H(C, N, F)$. All settings use $k = 8$, gate thresholds $(0.3, 0.7)$, and uncentered SVD with at most one flip pair per parent.

Setting	Gemma-3-12B			Qwen3-VL-8B		
	ℓ	α	$H(C, N, F)$	ℓ	α	$H(C, N, F)$
M3CoT exposure	12	1.0	61.6	4	1.0	78.3
M3CoT white balance	12	0.5	60.8	36	1.0	76.4
M3CoT JPEG	12	1.0	63.0	36	1.0	78.9
Robo2VLM exposure	12	0.5	73.4	36	1.0	82.4
Robo2VLM lens smudge	36	0.5	73.6	36	1.0	82.5
Robo2VLM motion blur	36	1.0	69.1	36	0.5	82.4
GMAI window level	12	0.5	70.0	12	0.5	70.1
GMAI defocus blur	36	1.0	70.9	36	1.0	63.7
GMAI rotation	36	0.5	64.2	36	0.5	70.7

C.2 STEERING IMPLEMENTATION

We extract the states used to estimate U at the final prompt position in unedited prefill, but apply the resulting directions only during decoding at the selected layer. The projection can be evaluated as $U(U^\top h)$ using two matrix-vector products, requiring $O(dk)$ arithmetic without materializing a $d \times d$ projection matrix. The evaluation harness computes both the unedited and steered passes at each decoding step, then uses the gate to select the output distribution.

Evidence visualization. Figure 4 uses a square-root rank axis and a linear cumulative-energy axis. Figure 5 uses Gaussian kernel density estimation with a shared bandwidth of 0.045 and reflection at the boundaries 0 and 1. Each token group is normalized to unit area within each backbone. A square-root density axis reveals overlap between the distributions; tick labels report density in the original units.

C.3 HYPERPARAMETER SELECTION

The six (ℓ, α) candidates are scored on validation triplets containing a clean image, one flip child, and one non-flip child. Parents follow a fixed 40-shard ordering; the first R shards are used to cover approximately 50 flip-bearing parents, with R adjusted to the setting’s flip density. Qwen3-VL defocus blur has 47 available validation flip pairs and approximately 24 selection triplets. Validation parents are disjoint from subspace-calibration and test parents. Varied images are stored as PNG at construction time and are regenerable exactly from the deterministic per-parent seed chain.

C.4 BASELINE ADAPTATIONS

We retain the original LEAD constants and official VTI strengths rather than retuning them for individual benchmark settings. The adaptations below address backbone compatibility. In contrast, FLIPDIR selects its method-specific intervention layer and strength on held-out validation data (Appendix C.3); test data are used only for final evaluation.

LEAD. We adapt the public Qwen2.5-VL implementation (Yang et al., 2025b; Huang et al., 2026) to each model’s embedding and token conventions. For Gemma-3, direct embedding inputs retain the model’s $\sqrt{d}$ scaling, termination recognizes its chat end token, and image-anchor lookup uses the model-specific image token. The original LEAD constants are unchanged; Table 3 reports this adapted implementation.

VTI. We use the official implementation and strengths. Qwen3-VL produces a variable visual-token grid, so we average VTI’s visual direction over tokens within each layer and apply it at all positions. The textual direction is unchanged. Setting $\alpha = \beta = 0$ reproduces the base model exactly.

C.5 DETERMINISTIC INFERENCE AND RELEASED ARTIFACTS

All open-weight experiments use greedy decoding, batch size one, fixed seeds, deterministic kernels, and TF32 disabled. Repeating one complete setting reproduced all 1,040 generations exactly. Exact reproduction uses the saved flip subspaces: GPU re-extraction can introduce numerical differences, and subspaces agreeing to 10^{-6} can still change H by several points at near-tied decoding steps.

Release contents. We will release model-specific benchmark instances, perturbation code and versioned configurations, inference and evaluation scripts, answer-normalization and flip-mining utilities, budget-matching tools, and extracted subspaces. Instance metadata includes the parent identifier, seed, perturbation parameters, prompt, raw output, and normalized answer.

D TRANSFER ACROSS DATASET AND PERTURBATION COMBINATIONS

Table 13 expands Figure 7 with all measured source-to-target transfers, including within-domain sources. Only U is transferred; the layer, strength, and gate thresholds use the target-specific configuration. Across the cross-domain combinations shown in Figure 7, at least 91% of target-matched H is retained.

Table 13: **Transfer results.** $H(C, N, F)$ using the source subspace and the target-specific layer, strength, and gate thresholds. The last column uses a target-matched subspace.

Model	Target setting	Source U	$H(C, N, F)$	Target-specific
Qwen3-VL-8B	Robo2VLM lens	M3CoT exposure	75.7	82.5
Qwen3-VL-8B	Robo2VLM lens	GMAI detector	76.3	82.5
Qwen3-VL-8B	M3CoT exposure	Robo2VLM lens	71.5	78.3
Qwen3-VL-8B	M3CoT exposure	GMAI detector	74.2	78.3
Qwen3-VL-8B	GMAI detector	M3CoT exposure	65.3	70.7
Qwen3-VL-8B	GMAI detector	Robo2VLM lens	69.7	70.7
Gemma-3-12B	M3CoT exposure	Robo2VLM lens	61.5	61.6
Gemma-3-12B	M3CoT exposure	GMAI detector	61.4	61.6
Gemma-3-12B	Robo2VLM lens	M3CoT exposure	69.4	73.6
Gemma-3-12B	Robo2VLM lens	GMAI detector	71.6	73.6
Gemma-3-12B	GMAI detector	M3CoT exposure	62.3	64.2
Gemma-3-12B	GMAI detector	Robo2VLM lens	60.5	64.2
Gemma-3-12B	M3CoT exposure	M3CoT JPEG	61.6	61.6
Gemma-3-12B	M3CoT exposure	Robo2VLM exposure	61.1	61.6
Gemma-3-12B	Robo2VLM lens	Robo2VLM exposure	71.1	73.6
Gemma-3-12B	GMAI detector	GMAI window	61.1	64.2
Qwen3-VL-8B	GMAI detector	GMAI window	60.9	70.7
Gemma-3-12B	Robo2VLM lens	Robo2VLM motion	71.6	73.6
Gemma-3-12B	GMAI detector	Robo2VLM motion	62.1	64.2
Qwen3-VL-8B	Robo2VLM lens	Robo2VLM motion	80.9	82.5
Qwen3-VL-8B	GMAI detector	Robo2VLM motion	70.2	70.7

E ABLATIONS AND SENSITIVITY

Figure 8 in Section 5.4 summarizes the steering-strength sweep, the local gate-threshold neighborhood, and the rank sweep around the shared configuration. The tables below give the full rank sweep, component ablations, and sensitivity to the number of subspace-estimation pairs.

E.1 FLIP-SUBSPACE RANK

Table 14 sweeps k on six settings, covering two non-medical perturbation families and one medical perturbation family per model, while holding all other components fixed at the selected configuration.

Table 14: **Rank sensitivity.** $H(C, N, F)$ with all other components fixed to Table 12. The shared rank $k = 8$ is shaded; row maxima are bold.

Model	Setting	$k = 2$	4	8	16	32
Qwen3-VL-8B	M3CoT exposure	71.0	74.3	78.3	73.2	72.6
	Robo2VLM lens smudge	82.2	79.6	82.5	79.1	77.8
	GMAI rotation	65.4	64.7	70.7	63.7	69.6
Gemma-3-12B	M3CoT exposure	59.2	58.6	61.6	59.2	60.2
	Robo2VLM lens smudge	77.1	70.4	73.6	71.4	70.2
	GMAI rotation	65.1	64.2	64.2	65.4	64.5

Spectral energy concentration. For Qwen3-VL-8B on Robo2VLM lens smudge, we analyze 255 clean-flip pairs ($d = 4096$) at layers 4, 12, and 36. With uncentered singular values $\sigma_{\ell,j}$, cumulative energy is $E_\ell(r) = \sum_{j=1}^r \sigma_{\ell,j}^2 / \sum_j \sigma_{\ell,j}^2$. The top 8 directions capture 48.6 to 53.7% of the energy, versus 4.7% for the Gaussian null; the top 16 capture 59.4 to 65.7%, versus 9.2%. Figure 4 plots these cumulative energies.

E.2 MARGIN-GATE ABLATION

Table 15 reports the preservation and recovery metrics underlying Table 6. Gate Off applies steering at every decoding step; the subspace, layer, and strength are held fixed within each comparison. The gate improves H in all six settings. Figure 9 visualizes the corresponding changes in each metric.

Table 15: **Margin-gate ablation.** $C/N/F/H$ (%) with and without the gate. $\Delta H = H_{\text{On}} - H_{\text{Off}}$, computed from the displayed values.

Setting	Gate On				Gate Off				ΔH
	C	N	F	H	C	N	F	H	
<i>Gemma-3-12B</i>									
M3CoT exposure	70.2	69.1	50.0	61.6	59.5	63.8	45.0	54.8	+6.8
Robo2VLM lens smudge	87.8	86.4	56.2	73.6	81.7	83.4	51.6	68.8	+4.8
GMAI rotation	84.5	83.2	43.7	64.2	84.0	78.7	37.3	58.3	+5.9
<i>Qwen3-VL-8B</i>									
M3CoT exposure	87.4	87.9	64.5	78.3	84.1	85.3	62.3	75.6	+2.7
Robo2VLM lens smudge	94.1	94.0	66.3	82.5	92.3	91.9	52.3	73.5	+9.0
GMAI rotation	80.2	81.8	56.4	70.7	79.2	81.8	53.6	69.2	+1.5




Figure 9: **Performance gains from margin gating.** Changes in clean preservation (C), non-flip preservation (N), flip recovery (F), and their harmonic mean (H), measured in percentage points as Gate On minus Gate Off. All panels share the same scale; values are computed from Table 15. The subspace, layer, and strength are fixed within each comparison. GMAI denotes GMAI-MMBench.

E.3 NUMBER OF EXAMPLES USED TO ESTIMATE THE FLIP SUBSPACE

Table 16 subsamples the examples used to estimate U , holding everything else at the selected configuration; a subsampled U uses the first N clean parents by sorted parent id (at most one flip-inducing perturbation per parent, uncentered SVD).

Table 16: **Subspace-estimation sample size.** $H(C, N, F)$ with other components fixed; parent counts are shown in parentheses.

Setting	50 pairs	subsampled	full pool
Gemma M3CoT exposure	60.2	56.8 (250)	61.6 (989)
Gemma Robo2VLM lens	67.0	72.2 (150)	73.6 (319)
Gemma GMAI detector	67.0	63.8 (100)	64.2 (192)
Qwen Robo2VLM lens	78.6	79.2 (150)	82.5 (255)

E.4 LIMITED PREDICTIVE VALUE OF STATIC DIAGNOSTICS

We examine whether simple diagnostics predict answer flips or recovery. Subspace alignment has little correlation with realized recovery ($\rho \approx 0$), and steering coordinates provide weak token-level discrimination (SNR 0.02 to 0.07; AUC 0.55). A parent’s flip probability is also nearly flat in its clean-run margin (0.48 to 0.60). These results limit the use of alignment, steering coordinates, or clean-run confidence as standalone predictors. FLIPDIR does not use U as a flip detector or clean-run confidence to predict future flips. It uses U as an intervention subspace and gates steering with the current unedited token margin during decoding of the perturbed input (Section 3). The diagnostics therefore distinguish static prediction from decoding-time intervention; they do not by themselves establish the benefit of dynamic gating. That benefit is evaluated directly by the controlled gate ablation (Table 6). We select (ℓ, α) using held-out validation performance rather than these proxies (Appendix C.3).

F COMPUTE, RUNTIME, AND GATE STATISTICS

F.1 COMPUTE ACCOUNTING

Offline direction construction. Direction estimation is performed once per backbone and perturbation setting, and the resulting subspace is reused across queries. Given existing clean and flip pairs, activation extraction requires two prefill-only forwards per pair; the candidate layers are read from the same forwards. The 4,512 Gemma pairs and 3,609 Qwen pairs in Table 10 therefore require 9,024 and 7,218 prefills, respectively. SVD is then applied to the activation differences, and

the retained basis stores dk coefficients. These extraction counts exclude pair search and validation-based parameter selection. Pair-search costs are reported separately in Table 8; validation selects the steering layer and strength before deployment. Thus the method incurs an initial construction cost in exchange for a reusable intervention, without updating model weights.

Decoding FLOPs. We count a multiply-add as two FLOPs and approximate the dominant linear operations in the decoder blocks and output head. Comparisons fix the input and output lengths and exclude prefill, image encoding, and the context-length-dependent attention products. Let F_b be the per-token cost of one decoder block, F_h the output-head cost, and L the number of blocks. Standard decoding costs $F_0 = LF_b + F_h$. Under conditional execution, an open gate triggers recomputation only above the intervention layer ℓ , including the output head. The relative cost is

$$\frac{F_{\text{FlipDir}}}{F_0} \approx 1 + p_{\text{open}} \frac{(L - \ell)F_b + F_h + 4dk}{LF_b + F_h}. \quad (6)$$

The projection $U(U^\top h)$ costs approximately $4dk$ FLOPs; gate arithmetic is negligible at this scale. The output head is recomputed even when $\ell = L$, so intervention at the final layer does not have zero cost.

For Gemma-3-12B, $L = 48$, $d = 3840$, $F_b \approx 4.48 \times 10^8$, and $F_h \approx 2.01 \times 10^9$ FLOPs. For Qwen3-VL-8B, $L = 36$, $d = 4096$, $F_b \approx 3.86 \times 10^8$, and $F_h \approx 1.24 \times 10^9$. We combine the layers in Table 12 with the gate-open fractions replayed on stored unedited traces in Table 17. Estimated multipliers range from 1.068 to 1.185 for Gemma and 1.011 to 1.176 for Qwen, across seven and eight settings with available traces, respectively. These replay fractions are proxies for conditional execution, rather than measurements from a deployed conditional implementation.

LEAD and VTI use one decoder forward per token plus small gating or additive-edit operations, giving approximately $1\times$ decoding FLOPs. VCD uses two forwards, giving approximately $2\times$. Under the same accounting, conditional FLIPDIR uses about 41 to 50% fewer decoding FLOPs than VCD. This comparison concerns decoding arithmetic rather than end-to-end latency. Our current evaluation harness computes both the unedited and steered paths at each step, so Table 7 describes the conditional execution of the method rather than measured FLOPs of that harness. Section F.3 reports its measured wall-clock runtime separately.

F.2 GATE ACTIVATION STATISTICS

Gate open rates replay the selected thresholds (0.3, 0.7) with hysteresis over stored unedited traces: twelve parents per setting across shards, including clean, flip, and non-flip generations. N/A marks settings without stored traces.

Table 17: **Runtime and gate statistics.** Time is measured in seconds per parent under the shared serving regime. Gate statistics are replayed on stored unedited traces; entries are counted per 100 tokens. N/A denotes unavailable measurements.

Setting	Gate		Seconds per parent			
	Open %	Entries	FLIPDIR	VCD	LEAD	VTI
Qwen M3CoT exposure	15.5	9.1	254.1	260.5	92.4	252.2
Qwen M3CoT wb	14.5	8.5	231.3	265.4	135.1	227.2
Qwen M3CoT JPEG	12.9	8.0	232.7	258.7	134.6	241.4
Qwen Robo2VLM exp.	17.3	9.6	121.8	109.0	71.9	44.7
Qwen Robo2VLM lens	N/A	N/A	71.8	96.0	49.2	43.7
Qwen Robo2VLM motion	18.8	10.2	112.7	120.4	70.8	42.1
Qwen GMAI window	25.4	13.0	85.7	209.8	113.8	88.6
Qwen GMAI defocus	24.1	12.2	74.5	209.7	114.2	80.5
Qwen GMAI rotation	24.0	12.7	76.8	211.2	115.7	82.8
Gemma M3CoT exposure	13.3	6.5	175.3	168.7	N/A	133.8
Gemma M3CoT wb	11.0	5.9	171.1	172.5	N/A	121.3
Gemma M3CoT JPEG	N/A	N/A	176.3	169.2	96.6	124.3
Gemma Robo2VLM exp.	22.3	11.6	48.0	44.1	N/A	40.8
Gemma Robo2VLM lens	N/A	N/A	36.4	35.5	17.7	38.3
Gemma Robo2VLM motion	24.4	10.8	22.3	44.0	N/A	53.2
Gemma GMAI window	24.0	10.9	53.6	97.5	N/A	89.1
Gemma GMAI defocus	21.7	11.8	49.2	92.6	N/A	74.5
Gemma GMAI rotation	25.8	12.9	54.4	101.5	N/A	91.0

F.3 RUNTIME

Wall-clock time in Table 17 is averaged over the main-table runs under a shared serving regime with four to five concurrent processes per GPU. The current FLIPDIR harness executes both the original and steered passes at each decoding step. Across the six medical settings, its measured runtime is 1.8 to $2.8\times$ faster than VCD while achieving higher H . LEAD uses a single forward pass.

G QUALITATIVE EXAMPLES

G.1 ANSWER-FLIP RECOVERY

Figure 10 illustrates Gemma-3-12B on Robo2VLM exposure shift with $\ell = 12$ and $\alpha = 0.5$. The clean answer is (B), matching the ground truth. Exposure shift changes the unedited answer to (A); FLIPDIR restores (B) on the same perturbed image. The displayed reasoning excerpts show the corresponding change from repeating the completed phase to predicting the next action.



Clean image (model answers B)



Exposure shift (-EV)

[Question] The robot is tasked to moving away with open gripper after releasing the can. After moving away with open gripper after releasing the can, what will be the robot's NEXT action phase?

- (A) Moving away with open gripper after releasing the can **(B) Approaching the can with open gripper**
 (C) Firmly grasping the can (D) Releasing the can by opening gripper (E) Closing gripper to grasp the can

Varied image, no steering → (A)

The question states that the robot has already “Moving away with open gripper after releasing the can”. The question asks for the *next* action. Since the can has already been released, the next logical action would be to move away with the open gripper.

FinalAnswer: (A)

Varied image, FLIPDIR → (B)

The question states that the robot has already “Moving away with open gripper after releasing the can”. This means the can has been released and the robot is moving away. The next logical action would be to approach another object or location.

Therefore, the next action phase is approaching the can with an open gripper.

FinalAnswer: (B)

Figure 10: A repaired flip on Gemma-3-12B Robo2VLM exposure shift. The unedited pass on the varied image restates the completed phase and flips to (A); with FLIPDIR the reasoning advances to the cycle's next phase and the answer returns to (B), the clean answer and the ground truth. Colored spans mark where the trajectories diverge.

G.2 PERTURBATION INTENSITY EXAMPLES

Figures 11, 12, and 13 provide representative examples from the fixed perturbation bands. Additional medical-specific design details are provided in Appendix A.5.




Figure 11: M3CoT variation intensities on the example of Figure 2, rendered from the fixed bands. Rows: exposure shift at the EV band endpoints (+0.45, +0.65), white-balance shift at +1000 K and +3000 K, and the mildest ($q = 95, 85$) and strongest ($q = 85, 72$) double-JPEG draws.




Figure 12: Robo2VLM lens-smudge example from ANSWERFLIPBENCH. The perturbation composites wipe smears, droplets, dust, local blur, and mild haze while preserving the underlying scene.




Figure 13: GMAI-MMBench variation intensities on an endoscopy image. Rows: window-level shift at the joint contrast and gamma endpoints, modality-gated defocus blur (disk PSF for this lens modality), and detector rotation at 1.5° and 4.0° .

H RECOVERY–PRESERVATION TRADE-OFFS

We plot $P = (C + N)/2$ against flip recovery F for all 18 settings in Tables 2 and 3. Each point represents the reported configuration, without additional tuning or interpolation. A point is Pareto-nondominated if no compared configuration has at least as high P and F and is strictly higher in one. FLIPDIR is nondominated in every setting under the reported rounded values, although other methods can remain on the frontier. Thus these comparisons show trade-offs independently of H , not dominance at every preservation level.




Figure 14: **Recovery–preservation trade-offs for Qwen3-VL-8B.** All nine settings and five methods are shown on shared axes. Higher and further right indicate better recovery and preservation, respectively.




Figure 15: **Recovery–preservation trade-offs for Gemma-3-12B.** All nine settings, with the same methods and axes as Figure 14. Base is at (100, 0) in every setting.